

# Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>  
(<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>  
(<https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

## Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

# [1]. Reading Data

## [1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [6]: !pip install kaggle
from google.colab import files
files.upload()
```

Requirement already satisfied: kaggle in /usr/local/lib/python3.6/dist-packages (1.5.2)  
 Requirement already satisfied: urllib3<1.23.0,>=1.15 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.22)  
 Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.11.0)  
 Requirement already satisfied: certifi in /usr/local/lib/python3.6/dist-packages (from kaggle) (2018.11.29)  
 Requirement already satisfied: python-dateutil in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.5.3)  
 Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.18.4)  
 Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages (from kaggle) (4.28.1)  
 Requirement already satisfied: python-slugify in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.0.1)  
 Requirement already satisfied: idna<2.7,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (2.6)  
 Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (3.0.4)  
 Requirement already satisfied: Unidecode>=0.04.16 in /usr/local/lib/python3.6/dist-packages (from python-slugify->kaggle) (1.0.23)

no files selected

Upload widget is only available when the cell has been executed in the current browser session.  
 Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

```
Out[6]: {'kaggle.json': b'{"username":"rohitbohra2994","key":"162fd3c566c99c9e6047f8248ccf8a9e"}'}
```

```
In [9]: !mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/

!chmod 600 ~/.kaggle/kaggle.json

!kaggle datasets download -d snap/amazon-fine-food-reviews

!ls
```

401 - Unauthorized  
 kaggle.json sample\_data

```
In [10]: !unzip amazon-fine-food-reviews.zip
```

```
unzip: cannot find or open amazon-fine-food-reviews.zip, amazon-fin  
e-food-reviews.zip.zip or amazon-fine-food-reviews.zip.ZIP.
```

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```

In [2]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 5
# you can change the number to any other number based on your computing

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """)

# Give reviews with Score>3 a positive rating(1), and reviews with a score less than 3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (525814, 10)

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	1
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

```
In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()

(80668, 7)
```

```
Out[4]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out[5]:
```

	UserId	ProductId	ProfileName	Time	Score	Text	COUI
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	

```
In [6]: display['COUNT(*)'].sum()
```

```
Out[6]: 393063
```

## [2] Exploratory Data Analysis

### [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time"})
final.shape
```

```
Out[9]: (364173, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 69.25890143662969
```



```
In [11]: #sorting of data
final['Time'] = pd.to_datetime(final['Time'])
# Sort by time
final = final.sort_values(by='Time')

print(final.shape)
print(final['Score'].value_counts())

print(final['Text'][1])
```

```
(364173, 10)
```

```
1    307063
```

```
0     57110
```

```
Name: Score, dtype: int64
```

Product arrived labeled as Jumbo Salted Peanuts...the peanuts were actually small sized unsalted. Not sure if this was an error or if the vendor intended to represent the product as "Jumbo".

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [12]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[12]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	3
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	3

```
In [13]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [14]: #Before starting the next phase of preprocessing lets see the number o.
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(364171, 10)
```

```
Out[14]: 1      307061
0       57110
Name: Score, dtype: int64
```

```
In [15]: final.head()
```

```
Out[15]:
```

	<b>Id</b>	<b>ProductId</b>	<b>UserId</b>	<b>ProfileName</b>	<b>HelpfulnessNumerator</b>	<b>Helpfulr</b>
<b>138706</b>	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	
<b>138683</b>	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	
<b>417839</b>	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	
<b>346055</b>	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	
<b>417838</b>	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	

## [3] Preprocessing

## [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [16]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

=====

I finally ordered a couple products from this seller for myself(not as gifts) and I am really happy. This Jade Bonsai is really cool and it arrived fast and in perfect condition. It's in my living room and I get tons of compliments. It's already grown some too and the pot it came in is really nice, looks expensive! Much bigger than I thought it would be even. Thanks again!!

=====

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.

=====

I would prefer freshly made brown rice, but that takes a long time to make and isn't easy. This makes it convenient, and takes all the guess work out of making it. I generally have been buying frozen organic brown rice, but that takes up lots of freezer space. The fact that this is easy to store at room temperature is a big plus. I'll be buying more.

=====

```
In [17]: # remove urls from text python: https://stackoverflow.com/a/40823105/4
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

```
In [18]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-extract-text-from-a-specific-div
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

=====

I finally ordered a couple products from this seller for myself(not as gifts) and I am really happy. This Jade Bonsai is really cool and it arrived fast and in perfect condition. It's in my living room and I get tons of compliments. It's already grown some too and the pot it came in is really nice, looks expensive! Much bigger than I thought it would be even. Thanks again!!

=====

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.

=====

I would prefer freshly made brown rice, but that takes a long time to make and isn't easy. This makes it convenient, and takes all the guess work out of making it. I generally have been buying frozen organic brown rice, but that takes up lots of freezer space. The fact that this is easy to store at room temperature is a big plus. I'll be buying more.

```
In [19]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [20]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

I bought some of this tea when I was in Seattle and I have been dying to get more. It really is the best tea I have ever had. It is great hot or cold.

=====

```
In [21]: #remove words with numbers python: https://stackoverflow.com/a/1808237
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

this witty little book makes my son laugh at loud. i recite it in the car as we're driving along and he always can sing the refrain. he's learned about whales, India, drooping roses: i love all the new words this book introduces and the silliness of it all. this is a classic book i am willing to bet my son will STILL be able to recite from memory when he is in college

```
In [22]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub(r'[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

I bought some of this tea when I was in Seattle and I have been dying to get more It really is the best tea I have ever had It is great hot or cold

```
In [23]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in t

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ou',
                "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves',
                'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its',
                'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'thi',
                'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'ha',
                'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',
                'at', 'by', 'for', 'with', 'about', 'against', 'between', 'i',
                'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out',
                'then', 'once', 'here', 'there', 'when', 'where', 'why', 'ho',
                'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so',
                's', 't', 'can', 'will', 'just', 'don', "don't", 'should',
                've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'c',
                "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't",
                "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
                'won', "won't", 'wouldn', "wouldn't"])
```

```
In [24]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower())
    preprocessed_reviews.append(sentence.strip())

100%|██████████| 364171/364171 [03:08<00:00, 1930.37it/s]
```

```
In [25]: # Random sampling
data = final.head(100000)

data_preprocessed_reviews=preprocessed_reviews[0:100000]
```

```
In [26]: print("The size of sampled data is ",data.shape)
print("The size of sampled data is ",len(data_preprocessed_reviews))

The size of sampled data is (100000, 10)
The size of sampled data is 100000
```



## [4] Featurization

### [4.1] BAG OF WORDS

```
In [27]: # Splitting into Train and test
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(data_preprocessed_
```

```
In [28]: #Bag of words
count_vect = CountVectorizer(max_features=2000, min_df=20)

x_train = count_vect.fit_transform(x_train)

x_test = count_vect.transform(x_test)
```

```
In [29]: om sklearn.ensemble import RandomForestClassifier
om sklearn.model_selection import TimeSeriesSplit
om sklearn.model_selection import GridSearchCV
om yellowbrick.model_selection import ValidationCurve

se_learners = [100,200,300,400,500]
pth = (range(1,20,4))
ram_grid={"n_estimators":base_learners,"max_depth":depth}
cv = TimeSeriesSplit(n_splits=5)
    = RandomForestClassifier(max_features='sqrt', class_weight="balanced")
v = GridSearchCV(rf, param_grid,scoring='roc_auc',cv=tscv,n_jobs=-1,ver
v.fit(x_train, y_train)
_scores = gsv.cv_results_['mean_test_score']
int("Model with best parameters :\n",gsv.best_estimator_)
int("Best Score: %.2f%%"%(gsv.best_score_*100))
int("Best HyperParameter: ",gsv.best_params_)
timal_n_estimators = gsv.best_estimator_.n_estimators
timal_max_depth = gsv.best_estimator_.max_depth

alidation Curve
z = ValidationCurve(gsv.best_estimator_, param_name="n_estimators",para
z.fit(x_train, y_train)
z.poof()
```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent w  
orkers.

[Parallel(n\_jobs=-1)]: Done 42 tasks | elapsed: 1.2min

[Parallel(n\_jobs=-1)]: Done 125 out of 125 | elapsed: 8.2min finish


ed

Model with best parameters :

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=13, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=500, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
```

Best Score: 88.69%

Best HyperParameter: {'max\_depth': 13, 'n\_estimators': 500}

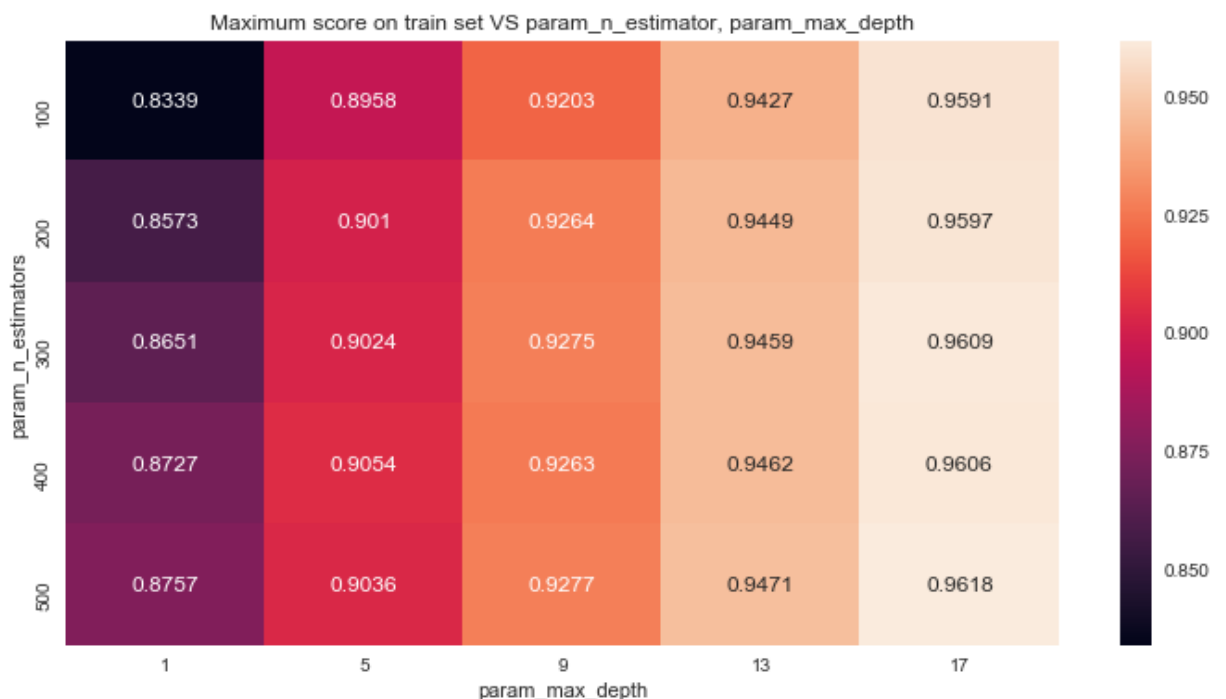


In [30]: gsv.best\_estimator\_


Out[30]: RandomForestClassifier(bootstrap=True, class\_weight='balanced',  
criterion='gini', max\_depth=13, max\_features='sqrt',  
max\_leaf\_nodes=None, min\_impurity\_decrease=0.0,  
min\_impurity\_split=None, min\_samples\_leaf=1,  
min\_samples\_split=2, min\_weight\_fraction\_leaf=0.0,  
n\_estimators=500, n\_jobs=None, oob\_score=False,  
random\_state=None, verbose=0, warm\_start=False)

```
In [31]: gsv.cv_results_
False,
        False],
        fill_value='?',
        dtype=object),
'param_n_estimators': masked_array(data=[100, 200, 300, 400, 500, 1
00, 200, 300, 400, 500, 100,
        200, 300, 400, 500, 100, 200, 300, 400, 500, 100
, 200,
        300, 400, 500],
        mask=[False, False, False, False, False, False, False,
False,
        False, False, False, False, False, False, False,
False,
        False, False, False, False, False, False, False,
False,
        False],
        fill_value='?',
        dtype=object),
'params': [{'max_depth': 1, 'n_estimators': 100},
{'max_depth': 1, 'n_estimators': 200},
```

```
In [34]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on train set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [32]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [35]: # This function plots the confusion, precision and recall matrices
def plot_confusion_matrix(x_test, y_pred):
    C = confusion_matrix(x_test, y_pred)

    A = ((C.T)/(C.sum(axis=1))).T
    B = (C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]

    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, fmt="d", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")

    plt.show()
```

In [36]: *ting Accuracy on Test data*

*rics*

```
sklearn.metrics import accuracy_score
sklearn.metrics import confusion_matrix
sklearn.metrics import precision_score
sklearn.metrics import recall_score
sklearn.metrics import f1_score
sklearn.metrics import roc_curve, auc
sklearn.metrics import roc_auc_score
```

*ting Accuracy on Test data*

```
print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)
rf = RandomForestClassifier(n_estimators=optimal_n_estimators, max_depth=optimal_max_depth)
rf.fit(x_train,y_train)
y_pred = rf.predict(x_test)
y_prob = rf.predict_proba(x_test)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
cm = confusion_matrix(y_test, y_pred)
```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 13


ROC\_AUC on test set: 0.898

Accuracy on test set: 83.733%

Precision on test set: 0.959

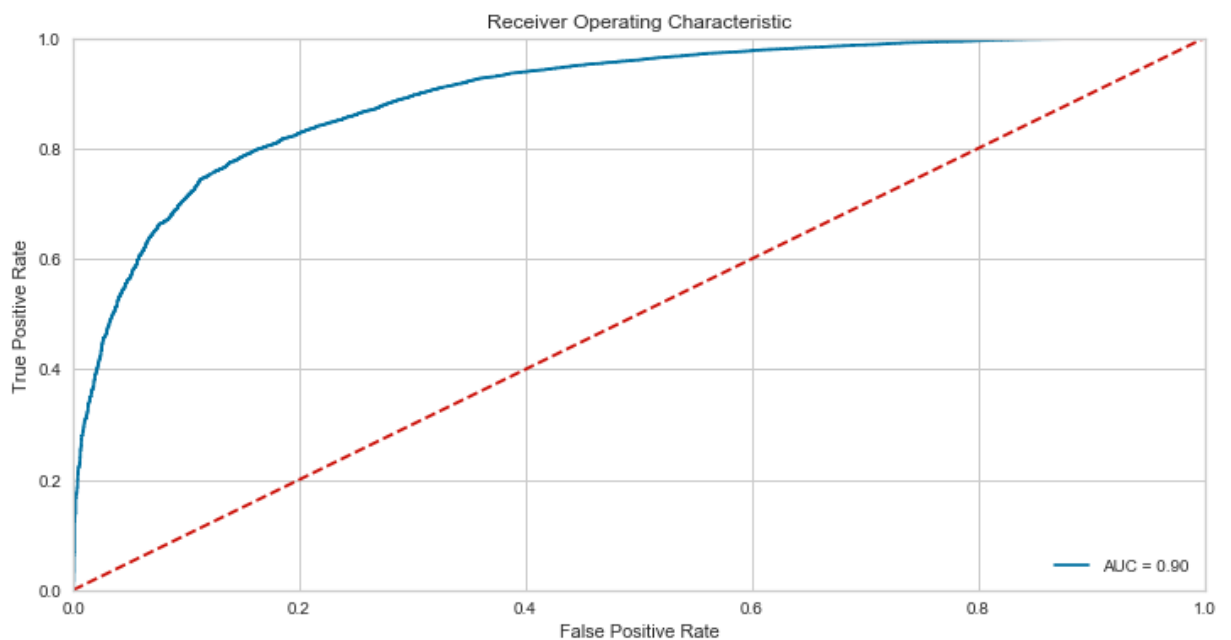
Recall on test set: 0.848

F1-Score on test set: 0.900



```
In [39]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [40]: #https://www.datacamp.com/community/tutorials/wordcloud-python
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt

feat_importances=rf.feature_importances_
feat_names=count_vect.get_feature_names()

# Sort feature importances in descending order
indices = np.argsort(feat_importances)[::-1][:25]
a=np.take(feat_names,indices)
def words(X):
    comment_words=' '
    for words in X:
        comment_words = comment_words + words + ' '
    return comment_words
a=words(a)

#Word Cloud
wc = WordCloud(max_font_size=30, max_words=100, background_color="white")
wc.generate(a)
plt.figure(figsize=[15,10])
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```



```
In [41]: from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
learning_rate = [0.01,0.05,0.1,0.2,0.3]

param_grid = {'n_estimators': base_learners, 'max_depth':depth, 'learning_rate':learning_rate}
cv = TimeSeriesSplit(n_splits=5)
```



```

gb = GradientBoostingClassifier(max_features='sqrt')
cv = GridSearchCV(gb, param_grid, scoring='roc_auc', cv=tscv, n_jobs=-1, verbose=1)
cv.fit(x_train, y_train)
scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n", gsv.best_estimator_)
print("Best HyperParameter: ", gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth
optimal_learning_rate = gsv.best_estimator_.learning_rate

Validation Curve
z = ValidationCurve(gsv.best_estimator_, param_name="n_estimators", param_range=(1, 1000))
z.fit(x_train, y_train)
z.poof()

```

Fitting 5 folds for each of 125 candidates, totalling 625 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 42 tasks | elapsed: 2.3min

[Parallel(n\_jobs=-1)]: Done 192 tasks | elapsed: 23.1min

/Users/rohitbohra/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process\_executor.py:706: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning

/Users/rohitbohra/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process\_executor.py:706: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning

/Users/rohitbohra/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process\_executor.py:706: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning

[Parallel(n\_jobs=-1)]: Done 442 tasks | elapsed: 57.6min

[Parallel(n\_jobs=-1)]: Done 625 out of 625 | elapsed: 86.9min finished

Model with best parameters :


```

GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=9,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=400,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=1.0, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)

```

Best HyperParameter: {'learning\_rate': 0.1, 'max\_depth': 9, 'n\_estimators': 400}

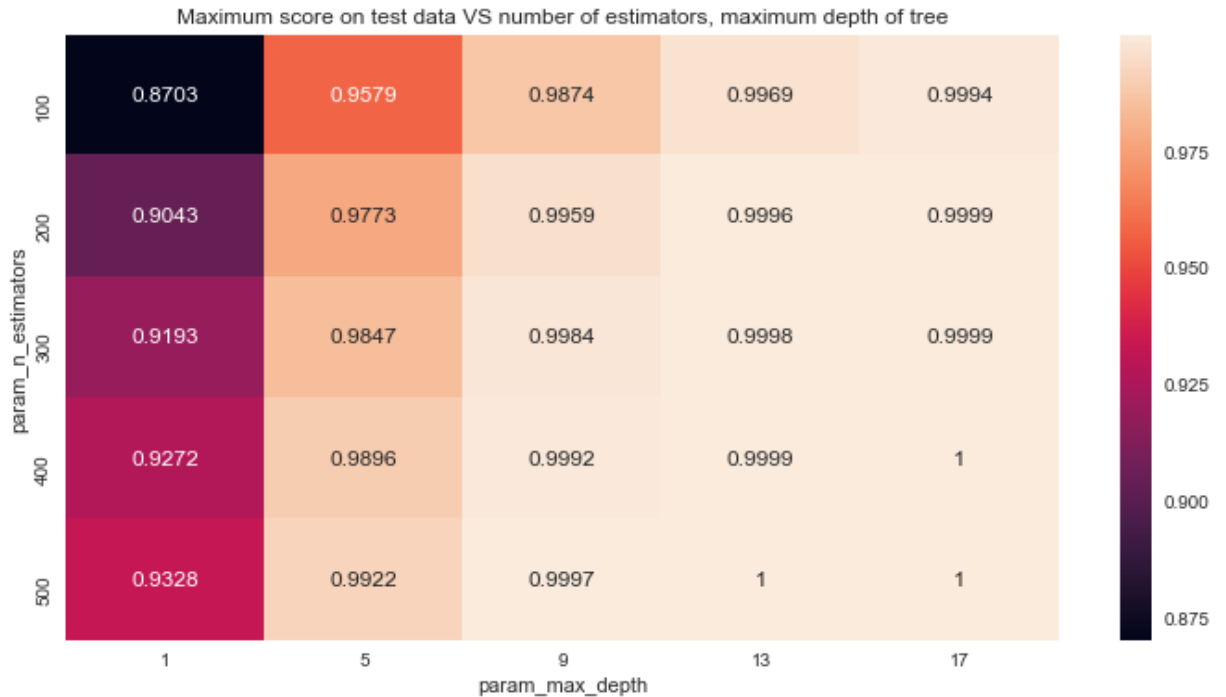
Best Accuracy: 92.04%



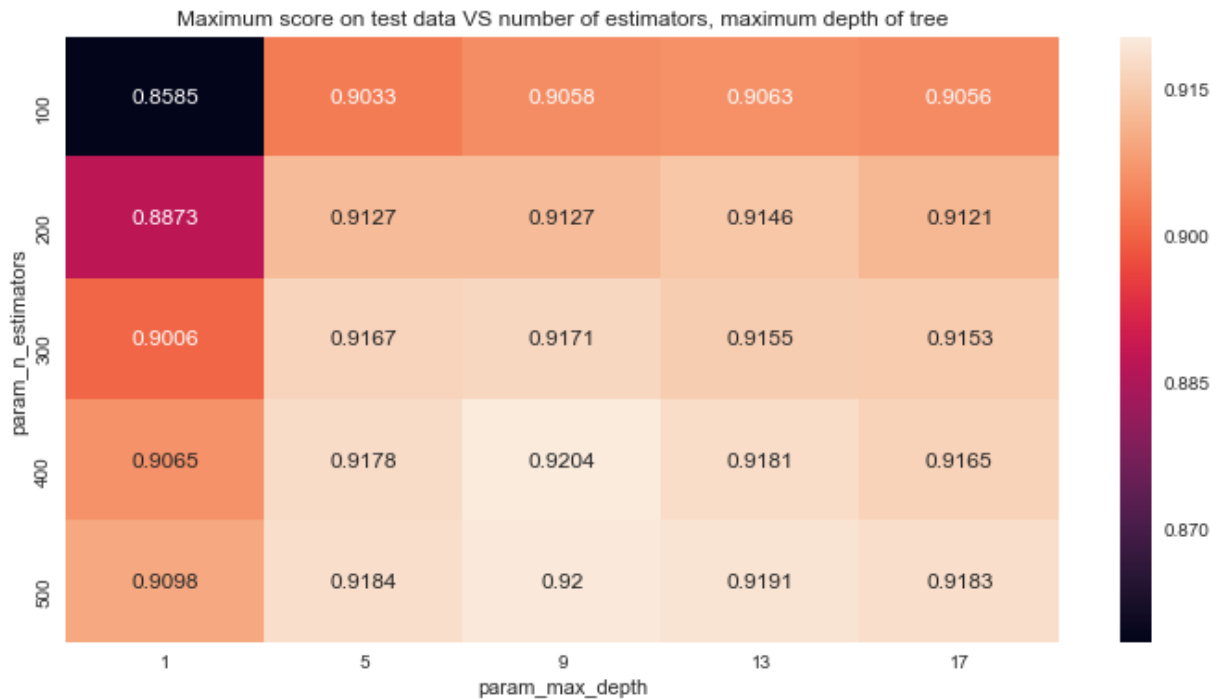
```
In [42]: gsv.best_estimator_
```

```
Out[42]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
learning_rate=0.1, loss='deviance', max_depth=9,
max_features='sqrt', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=400,
n_iter_no_change=None, presort='auto', random_state=None,
subsample=1.0, tol=0.0001, validation_fraction=0.1,
verbose=0, warm_start=False)
```

```
In [43]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [44]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



In [45]: *Testing Accuracy on Test data*

*metrics*

```

n sklearn.metrics import accuracy_score
n sklearn.metrics import confusion_matrix
n sklearn.metrics import precision_score
n sklearn.metrics import recall_score
n sklearn.metrics import f1_score
n sklearn.metrics import roc_curve, auc
n sklearn.metrics import roc_auc_score

```

*Testing Accuracy on Test data*

```

nt("The optimal value of n_estimators is : ",optimal_n_estimators)
nt("The optimal value of max_depth is : ",optimal_max_depth)

= GradientBoostingClassifier(n_estimators=optimal_n_estimators, max_dep
fit(x_train,y_train)
red = gb.predict(x_test)
rob = gb.predict_proba(x_test)

nt("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
nt("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100)
nt("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
nt("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
nt("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
t_confusion_matrix(y_test, y_pred)

```

The optimal value of n\_estimators is : 400

The optimal value of max\_depth is : 9


ROC\_AUC on test set: 0.932

Accuracy on test set: 91.210%

Precision on test set: 0.921


Recall on test set: 0.983

F1-Score on test set: 0.951



```
In [46]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [47]: #https://www.datacamp.com/community/tutorials/wordcloud-python
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt

importances=gb.feature_importances_
feat_names=count_vect.get_feature_names()

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1][:25]
a=np.take(feat_names,indices)
def words(X):
    comment_words=' '
    for words in X:
        comment_words = comment_words + words + ' '
    return comment_words
a=words(a)

#Word Cloud
wc = WordCloud(max_font_size=50, max_words=100, background_color="white")
wc.generate(a)
plt.figure(figsize=[20,10])
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```



## [4.3] TF-IDF

```
In [48]: from sklearn import preprocessing
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

# Splitting into Train and test
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(data_preprocessed_
tfidf = TfidfVectorizer(ngram_range=(1,2), max_features=2000, min_df=20)

#preparing the train data
x_train = tfidf.fit_transform(x_train)
#Normalize Data
x_train = preprocessing.normalize(x_train)
print("Train Data Size: ",x_train.shape)
print("the type of count vectorizer for train data is ",type(x_train))
print("the shape of train data is ",x_train.get_shape())
print("the number of unique words including both unigrams and bigrams :")
print("some sample features(unique words) ",tfidf.get_feature_names()[0])

#preparing the test data
x_test = tfidf.transform(x_test)
#Normalize Data
x_test = preprocessing.normalize(x_test)
print("Test Data Size: ",x_test.shape)
print("the type of count vectorizer for test data is ",type(x_test))
print("the shape of test data is ",x_test.get_shape())
print("the number of unique words including both unigrams and bigrams :")
```

```
Train Data Size: (70000, 2000)
the type of count vectorizer for train data is <class 'scipy.sparse
.csr.csr_matrix'>
the shape of train data is (70000, 2000)
the number of unique words including both unigrams and bigrams for t
rain data is 2000
some sample features(unique words) ['able', 'able find', 'absolute'
, 'absolutely', 'absolutely delicious', 'absolutely love', 'accordin
g', 'acid', 'across', 'actual']
Test Data Size: (30000, 2000)
the type of count vectorizer for test data is <class 'scipy.sparse.
csr.csr_matrix'>
the shape of test data is (30000, 2000)
the number of unique words including both unigrams and bigrams for t
est data is 2000
```

```
In [49]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import GridSearchCV
from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
```



```

param_grid={"n_estimators":base_learners,"max_depth":depth}
tscv = TimeSeriesSplit(n_splits=5)
rf = RandomForestClassifier(max_features='sqrt', class_weight="balanced")
gsv = GridSearchCV(rf, param_grid,scoring='roc_auc',cv=tscv,n_jobs=-1,
gsv.fit(x_train, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n",gsv.best_estimator_)
print("Best Score: %.2f%%"%(gsv.best_score_*100))
print("Best HyperParameter: ",gsv.best_params_)
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth

#Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators",p
viz.fit(x_train, y_train)
viz.poof()

```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent w
orkers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 125 out of 125 | elapsed: 10.4min finish
ed

```

Model with best parameters :

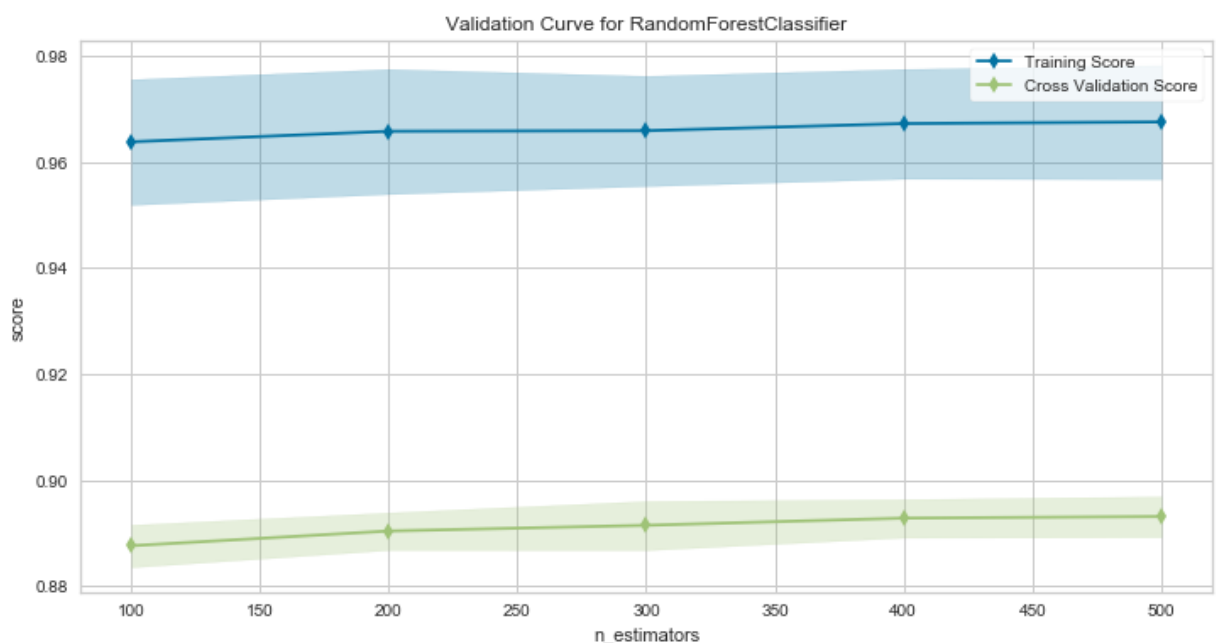
```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=17, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=500, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)


```

Best Score: 89.29%


Best HyperParameter: {'max\_depth': 17, 'n\_estimators': 500}



```
In [50]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [51]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



In [52]: *Testing Accuracy on Test data*

*Metrics*

```

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score

```

*Testing Accuracy on Test data*

```

print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)
f = RandomForestClassifier(n_estimators=optimal_n_estimators, max_depth=
f.fit(x_train,y_train)
_pred = rf.predict(x_test)
_prob = rf.predict_proba(x_test)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*10
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
plot_confusion_matrix(y_test, y_pred)

```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 17


ROC\_AUC on test set: 0.904

Accuracy on test set: 85.827%

Precision on test set: 0.955


Recall on test set: 0.877

F1-Score on test set: 0.914



```
In [53]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [54]: #https://www.datacamp.com/community/tutorials/wordcloud-python
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt

feat_importances=rf.feature_importances_
feat_names=count_vect.get_feature_names()

# Sort feature importances in descending order
indices = np.argsort(feat_importances)[::-1][:25]
a=np.take(feat_names,indices)
def words(X):
    comment_words=' '
    for words in X:
        comment_words = comment_words + words + ' '
    return comment_words
a=words(a)

#Word Cloud
wc = WordCloud(max_font_size=30, max_words=100, background_color="white")
wc.generate(a)
plt.figure(figsize=[15,10])
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```



```
In [55]: from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
Learning_rate = [0.01,0.05,0.1,0.2,0.3]

param_grid = {'n_estimators': base_learners, 'max_depth':depth, 'learning_rate': Learning_rate}
tscv = TimeSeriesSplit(n_splits=5)
gb = GradientBoostingClassifier(max_features='sqrt')
gsv = GridSearchCV(gb, param_grid,scoring='roc_auc',cv=tscv,n_jobs=-1,verbose=1)
gsv.fit(x_train, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n",gsv.best_estimator_)
print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth
optimal_learning_rate = gsv.best_estimator_.learning_rate
```

Fitting 5 folds for each of 125 candidates, totalling 625 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 42 tasks | elapsed: 2.6min

[Parallel(n\_jobs=-1)]: Done 192 tasks | elapsed: 25.9min

[Parallel(n\_jobs=-1)]: Done 442 tasks | elapsed: 61.3min

[Parallel(n\_jobs=-1)]: Done 625 out of 625 | elapsed: 89.2min finished


Model with best parameters :

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=9,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=500,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=1.0, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)
```


Best HyperParameter: {'learning\_rate': 0.1, 'max\_depth': 9, 'n\_estimators': 500}

Best Accuracy: 91.24%

```
In [56]: #Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators", p
viz.fit(x_train, y_train)
viz.poof()
```




```
In [57]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```





```
In [58]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```

In [59]: #Testing Accuracy on Test data
#Metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score

#Testing Accuracy on Test data
print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)

gb = GradientBoostingClassifier(n_estimators=optimal_n_estimators, max
gb.fit(x_train,y_train)
y_pred = gb.predict(x_test)
y_prob = gb.predict_proba(x_test)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1])
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred))
plot_confusion_matrix(y_test, y_pred)

```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 9


ROC\_AUC on test set: 0.933

Accuracy on test set: 91.343%

Precision on test set: 0.920

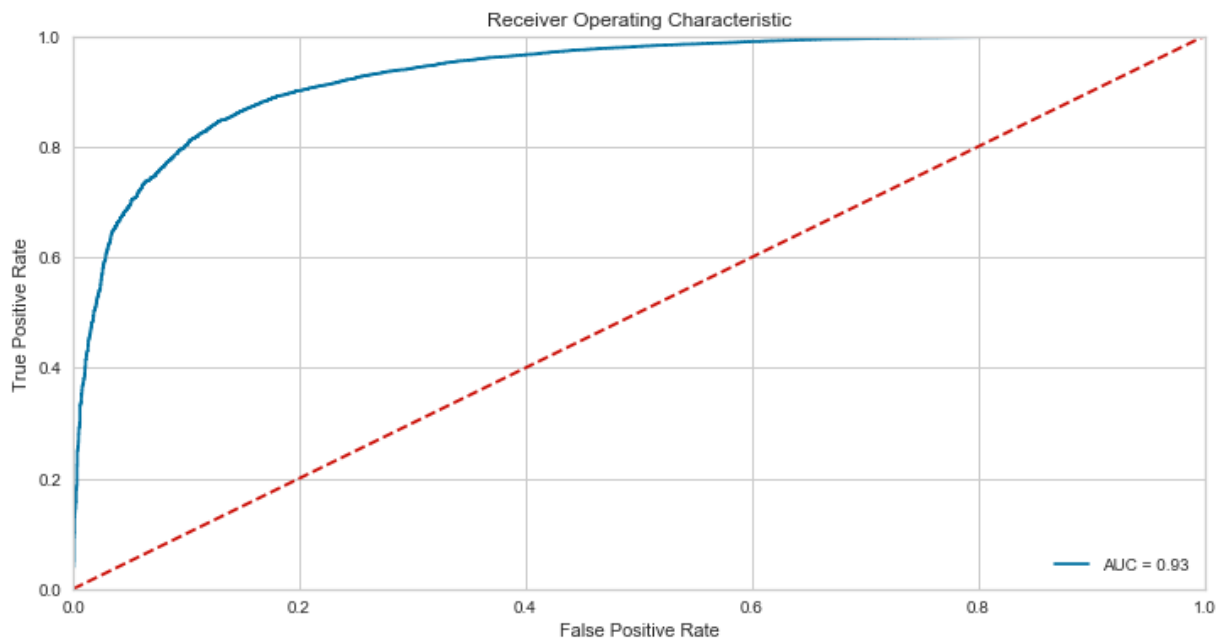
Recall on test set: 0.985

F1-Score on test set: 0.952



```
In [60]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```




```
In [61]: #https://www.datacamp.com/community/tutorials/wordcloud-python
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt

importances=gb.feature_importances_
feat_names=count_vect.get_feature_names()

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1][:25]
a=np.take(feat_names,indices)
def words(X):
    comment_words=' '
    for words in X:
        comment_words = comment_words + words + ' '
    return comment_words
a=words(a)

#Word Cloud
wc = WordCloud(max_font_size=50, max_words=100, background_color="white")
wc.generate(a)
plt.figure(figsize=[20,10])
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```



## [4.4] Word2Vec

```
In [62]: # Random sampling
data = final.head(50000)

data_preprocessed_reviews=preprocessed_reviews[0:50000]
print("The size of sampled data is ",data.shape)
print("The size of sampled data is ",len(data_preprocessed_reviews))
```

```
The size of sampled data is  (50000, 10)
The size of sampled data is  50000
```

```
In [63]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(data_preprocessed_
```

```
In [64]: # Train your own Word2Vec model using your train text corpus
i=0
list_of_train_sentence=[]
for sentence in x_train:
    list_of_train_sentence.append(sentence.split())
```

```
In [65]: # Train your own Word2Vec model using your test text corpus
i=0
list_of_test_sentence=[]
for sentence in x_test:
    list_of_test_sentence.append(sentence.split())
```

```

In [66]: # Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as va
# To use this code-snippet, download "GoogleNews-vectors-negative300.b
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/ed
# it's 1.9GB in size.

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W1
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occured atleast 5 times
    w2v_model=Word2Vec(list_of_train_sentence,min_count=5,size=50, worl
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vector
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_tra

[('wonderful', 0.8135449290275574), ('terrific', 0.812795102596283),
('excellent', 0.8054680824279785), ('fantastic', 0.804602861404419),
('awesome', 0.8004195690155029), ('good', 0.770828366279602), ('amaz
ing', 0.7531952857971191), ('perfect', 0.7331769466400146), ('fabulo
us', 0.7025632858276367), ('incredible', 0.6710256338119507)]
=====
[('greatest', 0.8081390857696533), ('best', 0.790654718875885), ('di
sgusting', 0.7568316459655762), ('tastiest', 0.7370363473892212), ('
closest', 0.7214698791503906), ('sweetest', 0.6875219345092773), ('n
icest', 0.6820852160453796), ('coolest', 0.6784370541572571), ('expe
rienced', 0.6771931648254395), ('awful', 0.6766822338104248)]

```

```
In [67]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

```
number of words that occurred minimum 5 times 11561
sample words ['witty', 'little', 'book', 'makes', 'son', 'laugh', 'loud', 'car', 'driving', 'along', 'always', 'sing', 'learned', 'india', 'roses', 'love', 'new', 'words', 'classic', 'willing', 'bet', 'stall', 'able', 'memory', 'college', 'remember', 'seeing', 'show', 'television', 'years', 'ago', 'child', 'sister', 'later', 'bought', 'day', 'thirty', 'something', 'used', 'series', 'books', 'songs', 'student', 'teaching', 'turned', 'whole', 'school', 'purchasing', 'cd', 'children']
```

## [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

### [4.4.1.1] Avg W2v

```
In [68]: average Word2Vec
compute average word2vec for each review in Train data.
sent_train_vectors = []; # the avg-w2v for each sentence/review is stored
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_train_vectors.append(sent_vec)
print(len(sent_train_vectors))
print(len(sent_train_vectors[0]))
```

```
100%|██████████| 35000/35000 [01:27<00:00, 399.56it/s]
```

```
35000
```

```
50
```

```
In [69]: # average Word2Vec
# compute average word2vec for each review in Test data.
sent_test_vectors = []; # the avg-w2v for each sentence/review is stored
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, y
    cnt_words = 0; # num of words with a valid vector in the sentence/r
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_test_vectors.append(sent_vec)
print(len(sent_test_vectors))
print(len(sent_test_vectors[0]))
```

100%|██████████| 15000/15000 [00:39<00:00, 380.74it/s]

15000

50

```
In [70]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import GridSearchCV
from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
param_grid={"n_estimators":base_learners,"max_depth":depth}
tscv = TimeSeriesSplit(n_splits=5)
rf = RandomForestClassifier(max_features='sqrt', class_weight="balanced")
gsv = GridSearchCV(rf, param_grid,scoring='roc_auc',cv=tscv,n_jobs=-1,verbose=0)
gsv.fit(sent_train_vectors, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n",gsv.best_estimator_)
print("Best Score: %.2f%%"%(gsv.best_score_*100))
print("Best HyperParameter: ",gsv.best_params_)
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth

Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators",param_grid=param_grid)
viz.fit(sent_train_vectors, y_train)
viz.poof()
```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.



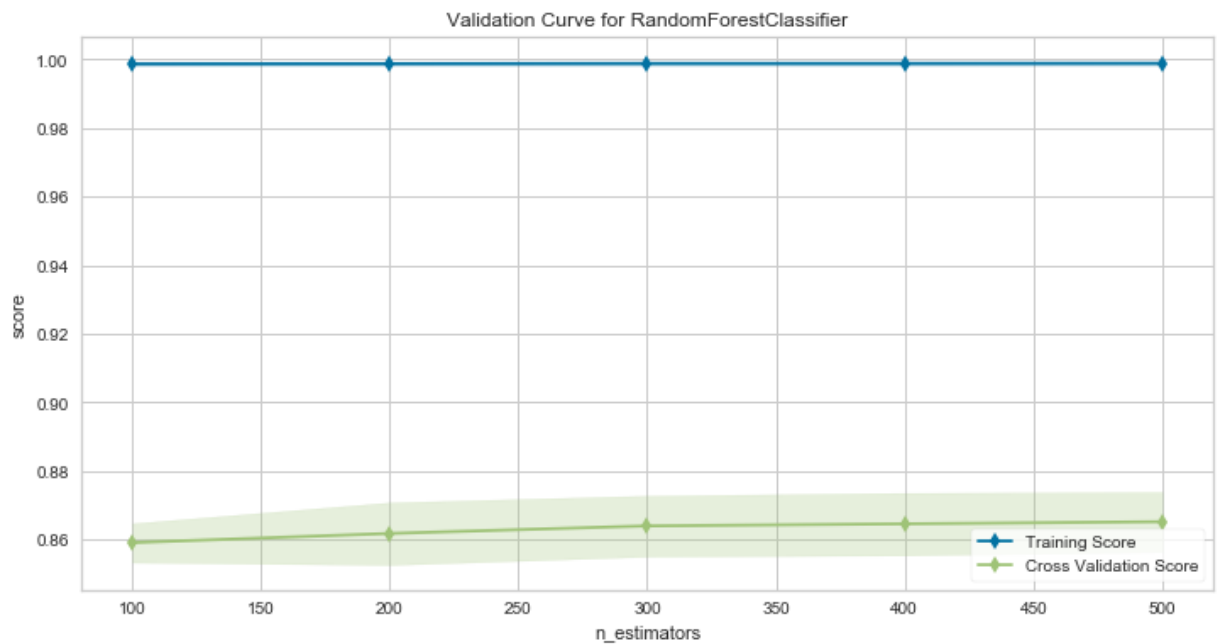
```
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 3.0min  
[Parallel(n_jobs=-1)]: Done 125 out of 125 | elapsed: 19.2min finished
```

Model with best parameters :

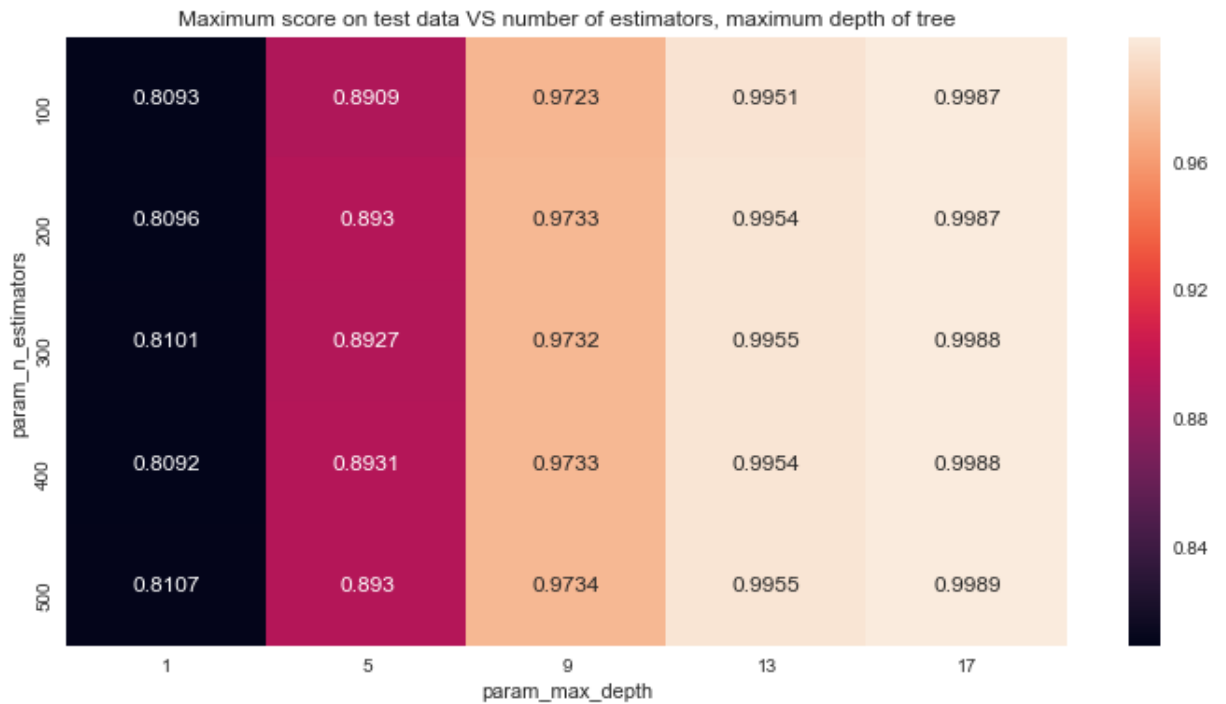
```
RandomForestClassifier(bootstrap=True, class_weight='balanced',  
                        criterion='gini', max_depth=17, max_features='sqrt',  
                        max_leaf_nodes=None, min_impurity_decrease=0.0,  
                        min_impurity_split=None, min_samples_leaf=1,  
                        min_samples_split=2, min_weight_fraction_leaf=0.0,  
                        n_estimators=500, n_jobs=None, oob_score=False,  
                        random_state=None, verbose=0, warm_start=False)
```

Best Score: 86.45%

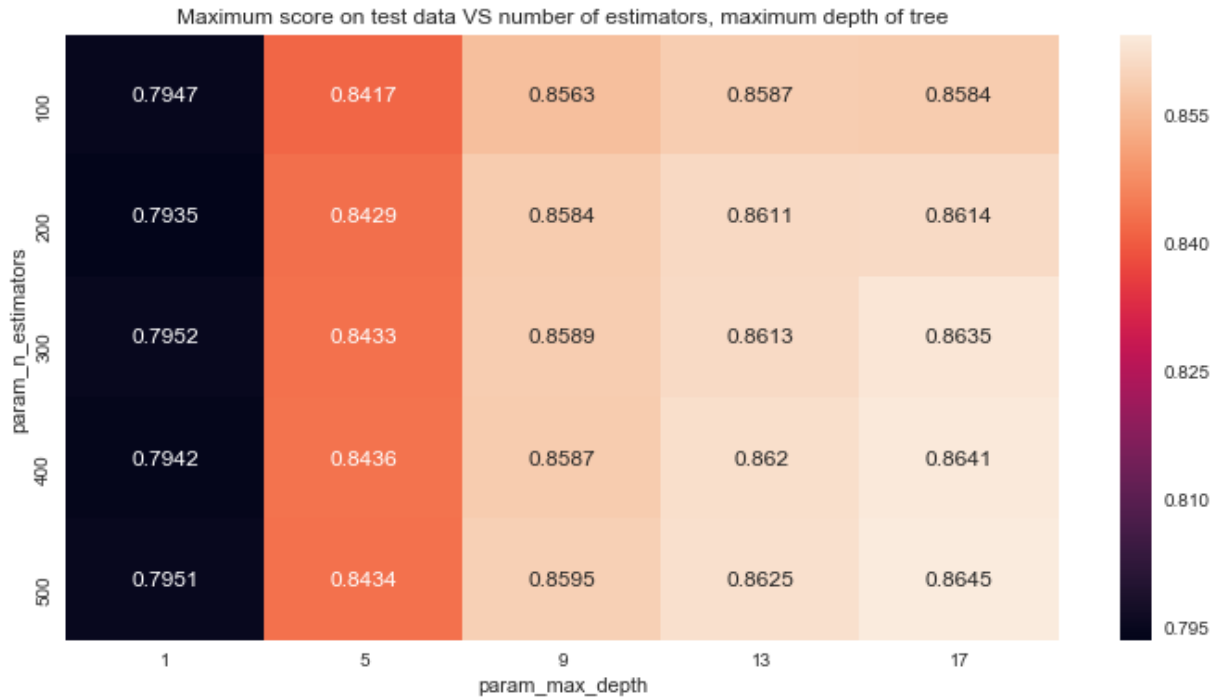
Best HyperParameter: {'max\_depth': 17, 'n\_estimators': 500}



```
In [71]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [72]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



In [73]: *Testing Accuracy on Test data*

*Metrics*

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score
```

*Testing Accuracy on Test data*

```
print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)
rf = RandomForestClassifier(n_estimators=optimal_n_estimators, max_depth=
    optimal_max_depth)
rf.fit(sent_train_vectors,y_train)
ypred = rf.predict(sent_test_vectors)
proba = rf.predict_proba(sent_test_vectors)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_proba[:,1]))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
cm = confusion_matrix(y_test, y_pred)
```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 17


ROC\_AUC on test set: 0.857

Accuracy on test set: 88.727%

Precision on test set: 0.911


Recall on test set: 0.967

F1-Score on test set: 0.938



```
In [74]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



```
In [75]: from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import TimeSeriesSplit

from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
Learning_rate = [0.01,0.05,0.1,0.2,0.3]

param_grid = {'n_estimators': base_learners, 'max_depth':depth, 'learning_rate': Learning_rate}
tscv = TimeSeriesSplit(n_splits=5)
gb = GradientBoostingClassifier(max_features='sqrt')
gsv = GridSearchCV(gb, param_grid, scoring='roc_auc', cv=tscv, n_jobs=-1, verbose=1)
gsv.fit(sent_train_vectors, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n", gsv.best_estimator_)
print("Best HyperParameter: ", gsv.best_params_)
print("Best Accuracy: %.2f%%" % (gsv.best_score_*100))
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth
optimal_learning_rate = gsv.best_estimator_.learning_rate
```

Fitting 5 folds for each of 125 candidates, totalling 625 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 42 tasks	elapsed: 3.4min
[Parallel(n_jobs=-1)]: Done 192 tasks	elapsed: 129.8min
[Parallel(n_jobs=-1)]: Done 442 tasks	elapsed: 241.8min
[Parallel(n_jobs=-1)]: Done 625 out of 625	elapsed: 280.9min finished


Model with best parameters :

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.05, loss='deviance', max_depth=5,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=500,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=1.0, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)
```

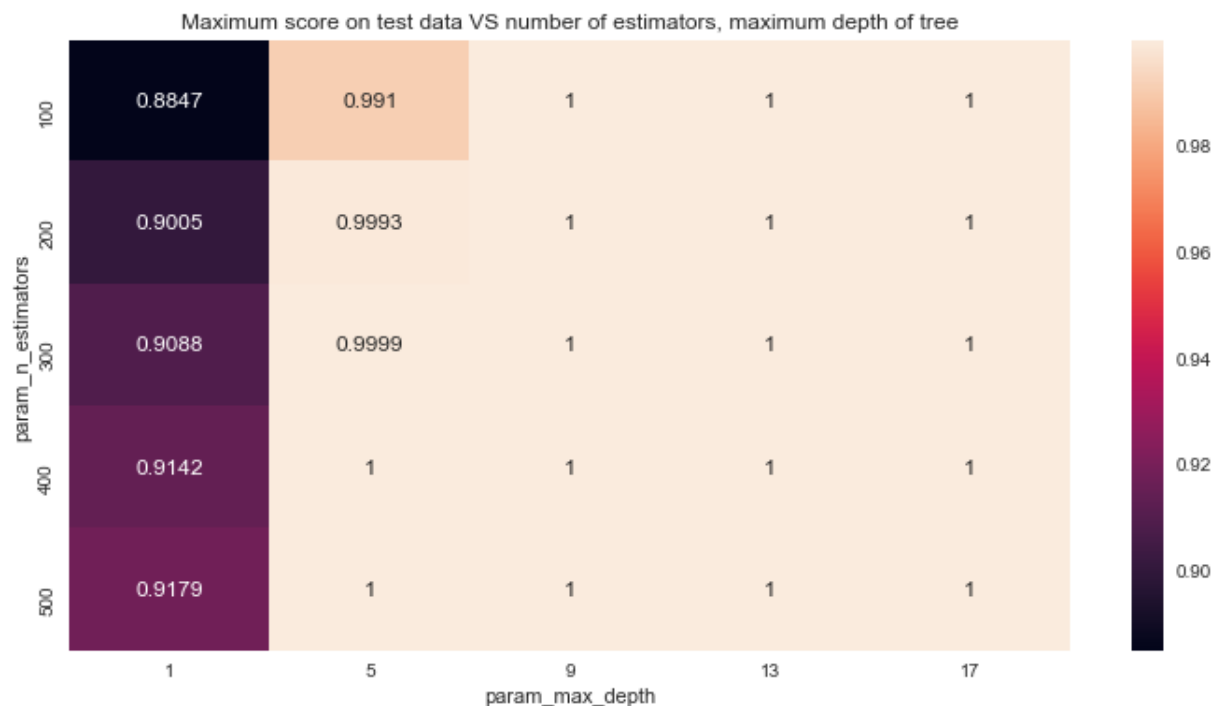
Best HyperParameter: {'learning\_rate': 0.05, 'max\_depth': 5, 'n\_estimators': 500}

Best Accuracy: 88.28%

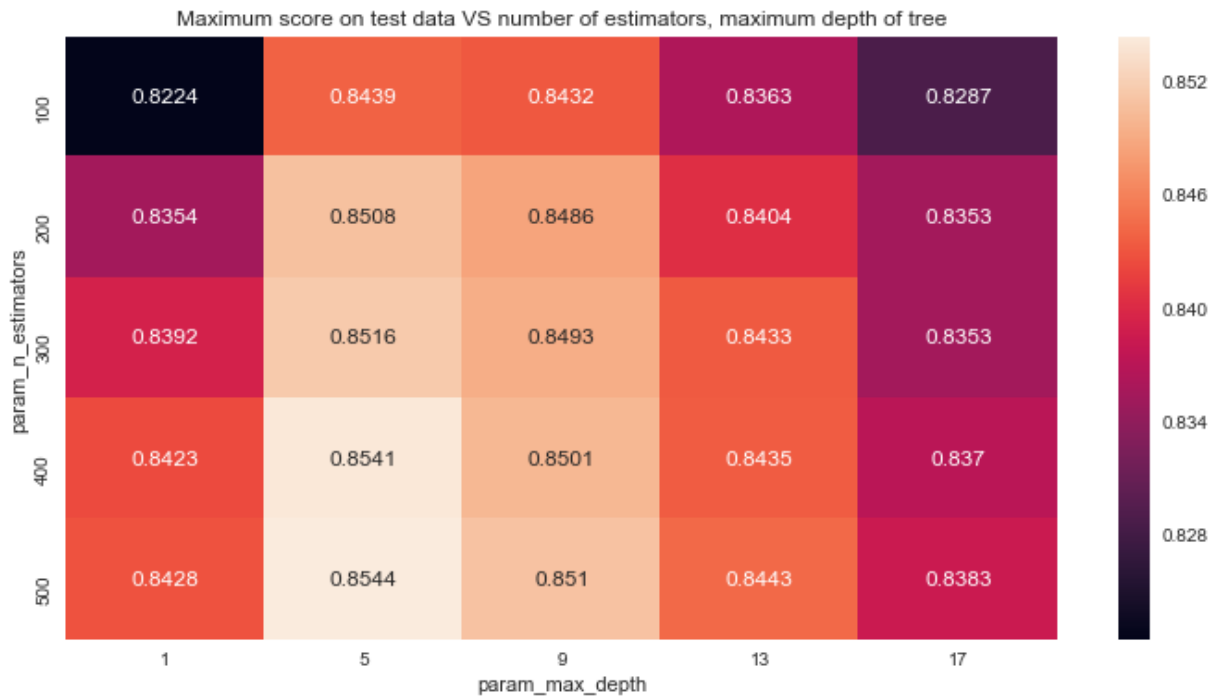
```
In [76]: #Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators", p
viz.fit(sent_train_vectors, y_train)
viz.poof()
```



```
In [78]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [95]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test data VS number of estimators, maximum depth of tree'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```





In [79]: *string Accuracy on Test data*  
*trics*

```
m sklearn.metrics import accuracy_score
m sklearn.metrics import confusion_matrix
m sklearn.metrics import precision_score
m sklearn.metrics import recall_score
m sklearn.metrics import f1_score
m sklearn.metrics import roc_curve, auc
m sklearn.metrics import roc_auc_score

string Accuracy on Test data
nt("The optimal value of n_estimators is : ",optimal_n_estimators)
nt("The optimal value of max_depth is : ",optimal_max_depth)

= GradientBoostingClassifier(n_estimators=optimal_n_estimators, max_dep
fit(sent_train_vectors,y_train)
red = gb.predict(sent_test_vectors)
rob = gb.predict_proba(sent_test_vectors)

nt("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
nt("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
nt("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
nt("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
nt("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
t_confusion_matrix(y_test, y_pred)
```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 5


ROC\_AUC on test set: 0.878

Accuracy on test set: 89.767%

Precision on test set: 0.911


Recall on test set: 0.980

F1-Score on test set: 0.944



```
In [80]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



#### [4.4.1.2] TFIDF weighted W2v

```
In [81]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer(ngram_range=(1,1))
tf_idf_matrix = model.fit_transform(data_preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [82]: # TF-IDF weighted Word2Vec on Train data
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and c

train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review
row=0;
for sent in tqdm(list_of_train_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    train_tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100%|██████████| 35000/35000 [18:34<00:00, 31.39it/s]

```
In [83]: # TF-IDF weighted Word2Vec on Test data

test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review
row=0;
for sent in tqdm(list_of_test_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    test_tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100%|██████████| 15000/15000 [08:18<00:00, 30.12it/s]

```
In [84]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import GridSearchCV
from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
param_grid={"n_estimators":base_learners,"max_depth":depth}
tscv = TimeSeriesSplit(n_splits=5)
rf = RandomForestClassifier(max_features='sqrt', class_weight="balanced")
gsv = GridSearchCV(rf, param_grid, scoring='roc_auc', cv=tscv, n_jobs=-1)
```

```

gsv.gridsearcher(11, param_grid, scoring = 'f1', cv = 5, n_jobs = -1)
gsv.fit(train_tfidf_sent_vectors, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n",gsv.best_estimator_)
print("Best Score: %.2f%%"%(gsv.best_score_*100))
print("Best HyperParameter: ",gsv.best_params_)
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth

#Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators", p
viz.fit(train_tfidf_sent_vectors, y_train)
viz.poof()

```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 42 tasks | elapsed: 3.0min  
 /Users/rohitbohra/anaconda3/lib/python3.6/site-packages/sklearn/externals/joblib/externals/loky/process\_executor.py:706: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.

"timeout or by a memory leak.", UserWarning

[Parallel(n\_jobs=-1)]: Done 125 out of 125 | elapsed: 19.4min finished

Model with best parameters :

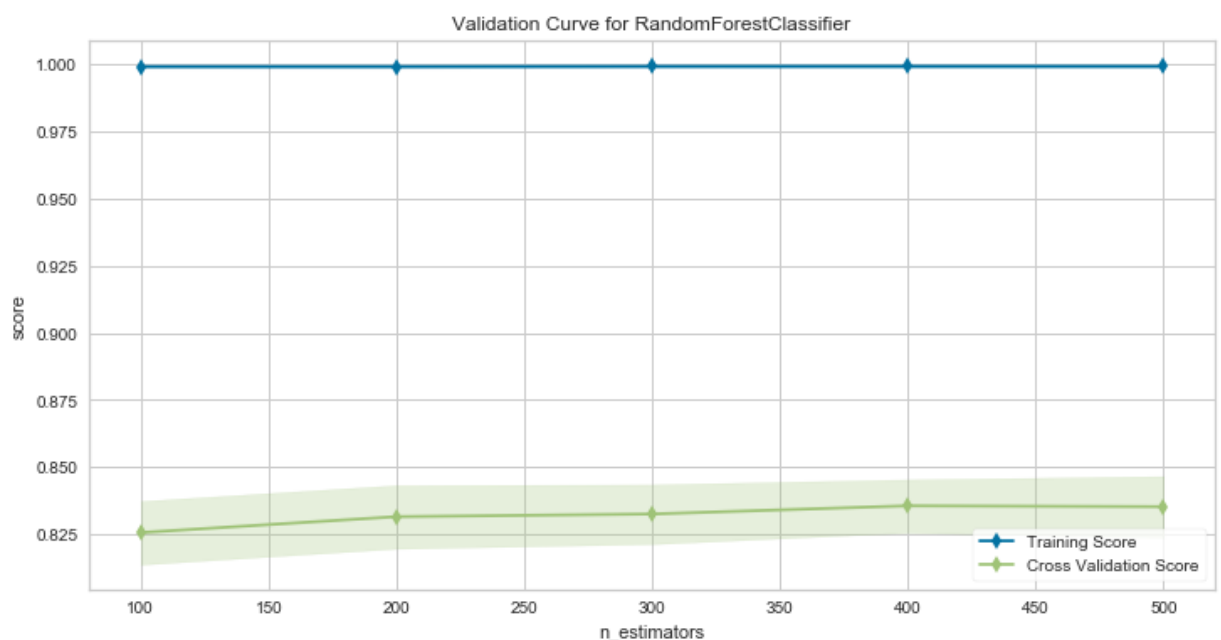
```

RandomForestClassifier(bootstrap=True, class_weight='balanced',
                        criterion='gini', max_depth=17, max_features='sqrt',
                        max_leaf_nodes=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=500, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)


```

Best Score: 83.49%


Best HyperParameter: {'max\_depth': 17, 'n\_estimators': 500}



```
In [85]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [86]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



In [87]: *Testing Accuracy on Test data*

*Metrics*

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score
```

*Testing Accuracy on Test data*

```
print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)
f = RandomForestClassifier(n_estimators=optimal_n_estimators, max_depth=optimal_max_depth)
f.fit(sent_train_vectors,y_train)
_pred = rf.predict(test_tfidf_sent_vectors)
_prob = rf.predict_proba(test_tfidf_sent_vectors)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
plot_confusion_matrix(y_test, y_pred)
```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 17


ROC\_AUC on test set: 0.830

Accuracy on test set: 88.800%

Precision on test set: 0.902


Recall on test set: 0.979

F1-Score on test set: 0.939



```
In [88]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```





```
In [89]: from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import TimeSeriesSplit

from yellowbrick.model_selection import ValidationCurve

base_learners = [100,200,300,400,500]
depth = (range(1,20,4))
Learning_rate = [0.01,0.05,0.1,0.2,0.3]

param_grid = {'n_estimators': base_learners, 'max_depth':depth, 'learning_rate': Learning_rate}
tscv = TimeSeriesSplit(n_splits=5)
gb = GradientBoostingClassifier(max_features='sqrt')
gsv = GridSearchCV(gb, param_grid, scoring='roc_auc', cv=tscv, n_jobs=-1)
gsv.fit(train_tfidf_sent_vectors, y_train)
cv_scores = gsv.cv_results_['mean_test_score']
print("Model with best parameters :\n", gsv.best_estimator_)
print("Best HyperParameter: ", gsv.best_params_)
print("Best Accuracy: %.2f%%" % (gsv.best_score_*100))
optimal_n_estimators = gsv.best_estimator_.n_estimators
optimal_max_depth = gsv.best_estimator_.max_depth
optimal_learning_rate = gsv.best_estimator_.learning_rate
```

Fitting 5 folds for each of 125 candidates, totalling 625 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 42 tasks	elapsed: 2.8min
[Parallel(n_jobs=-1)]: Done 192 tasks	elapsed: 129.8min
[Parallel(n_jobs=-1)]: Done 442 tasks	elapsed: 238.8min
[Parallel(n_jobs=-1)]: Done 625 out of 625	elapsed: 284.0min finished


Model with best parameters :

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.05, loss='deviance', max_depth=5,
                           max_features='sqrt', max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=500,
                           n_iter_no_change=None, presort='auto', random_state=None,
                           subsample=1.0, tol=0.0001, validation_fraction=0.1,
                           verbose=0, warm_start=False)
```


Best HyperParameter: {'learning\_rate': 0.05, 'max\_depth': 5, 'n\_estimators': 500}

Best Accuracy: 85.44%


```
In [90]: #Validation Curve
viz = ValidationCurve(gsv.best_estimator_, param_name="n_estimators", p
viz.fit(train_tfidf_sent_vectors, y_train)
viz.poof()
```



```
In [92]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_train_score, annot=True, fmt='.4g');
plt.title(title);
```



```
In [93]: df_gridsearch = pd.DataFrame(gsv.cv_results_)
max_scores = df_gridsearch.groupby(['param_n_estimators', 'param_max_depth'])
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12, 6)
title = 'Maximum score on test set VS param_n_estimator, param_max_depth'
fmt = 'png'
sns.heatmap(max_scores.mean_test_score, annot=True, fmt='.4g');
plt.title(title);
```



In [91]: *Testing Accuracy on Test data*

*Metrics*

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score
```

*Testing Accuracy on Test data*

```
print("The optimal value of n_estimators is : ",optimal_n_estimators)
print("The optimal value of max_depth is : ",optimal_max_depth)

gb = GradientBoostingClassifier(n_estimators=optimal_n_estimators, max_depth=optimal_max_depth)
gb.fit(train_tfidf_sent_vectors,y_train)
y_pred = gb.predict(test_tfidf_sent_vectors)
y_prob = gb.predict_proba(test_tfidf_sent_vectors)

print("ROC_AUC on test set: %0.3f" % roc_auc_score(y_test, y_prob[:,1]))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, y_pred)*100))
print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred)))
print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred)))
print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred)))
plot_confusion_matrix(y_test, y_pred)
```

The optimal value of n\_estimators is : 500

The optimal value of max\_depth is : 5


ROC\_AUC on test set: 0.855

Accuracy on test set: 89.160%

Precision on test set: 0.903


Recall on test set: 0.982

F1-Score on test set: 0.941



```
In [94]: preds = y_prob[:,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, preds)
roc_auc = metrics.auc(fpr, tpr)

# method 1: plt
import matplotlib.pyplot as plt
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



Ensemble Modeling								
Ensemble Technique		Featurization		Number of Estimator	Maximum Depth	ROC-AUC	F1-Score	Accuracy
Random Forest		BOW		500	13	0.898	0.9	83.733
Gradient Boosting DT		BOW		400	9	0.932	0.951	91.21
Random Forest		TF-IDF		500	17	0.904	0.914	85.827
Gradient Boosting DT		TF-IDF		500	9	0.933	0.952	91.343
Random Forest		Avg. W2V		500	17	0.857	0.938	88.727
Gradient Boosting DT		Avg-W2V		500	5	0.878	0.944	89.767
Random Forest		TF-IDF Avg. W2V		500	17	0.83	0.939	88.8
Gradient Boosting DT		TF-IDF Avg. W2V		500	5	0.855	0.941	89.16

# Conclusion

- 1) The Best ROC\_AUC score was of TF-IDF featurization and GBDT technique with maximum depth of 9 and number of estimator of 500.
- 2) The Best Accuracy score was of TF-IDF featurization and GBDT technique with maximum depth of 9 and number of estimator of 500.
- 3) GBDT technique takes alot of time to train and execute when compared to Decision Tree but has increased the model score by around 3-5% in BOW and TF-IDF featurization.

In [0]: